Source: https://codetrappers.com/ai-monitoring-observability/

Know what your agents actually did, and what it cost

Traces, token cost, retrieval logs and evaluation harnesses for AI systems already in production. Know what your agents did, what it cost, and when quality slipped.

An agent that quietly started failing three weeks ago is the normal failure mode, not a dramatic one. Observability is how you find out on the day it happens rather than from a customer, and how you answer the question an auditor will eventually ask.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

- 300+ migrations, zero data loss
- Core Web Vitals in the green
- 24/7 emergency cover

What you get

Inside AI Monitoring and Observability

Get in touch (https://codetrappers.com/contact/)

Trace every call from question to answer

What was asked, which tools ran with which arguments, what came back, and how long each step took. Without that, debugging an AI system is guesswork dressed up as intuition.

Treat cost as a first-class number

Tokens and spend per run, per feature and per customer, watched on the same dashboard as errors. AI bills do not creep, they step, and the step is usually a prompt change nobody flagged.

Record what was retrieved, not only what was said

When an answer is wrong, the useful question is almost always which passages the system pulled. Logging the retrieval alongside the response turns a complaint into a fixable bug.

Catch the quiet failures

Refusal rates climbing, retries increasing, latency drifting, a tool silently returning empty since a supplier changed a field. None of these page anybody. All of them are visible in the numbers before anyone notices.

Prove it still works after every change

A set of tasks with known good outcomes, run on every prompt edit, model upgrade

and deployment. Without one, every improvement is anecdotal and every regression is a surprise.

The service

What we build

Request and response tracing

End to end traces through your agent, its tools and its retrieval, kept where your team already looks rather than in another tool nobody opens.

Token and cost monitoring

Spend broken down by feature, model and tenant, with budgets and alerts so an unexpected bill arrives as a notification rather than an invoice.

Retrieval and citation logging

Which documents were used for which answer, retained long enough to investigate a complaint properly.

Agent action audit trails

Which tool ran, under whose credentials, what changed and what was refused. The record that makes an incident explainable rather than mysterious.

Evaluation harnesses and regression suites

Task sets with expected outcomes, scored automatically, so quality is a number in your pipeline rather than an impression after a demo.

Alerting and reporting

Thresholds on the things that actually matter, and periodic reports in a form a non-engineer can read and a compliance team can file.

Why CodeTrappers

What observability will not do

Three things worth being clear about before anyone buys a dashboard.

It does not stop hallucination

It shows you where and how often it happens, which is what lets you fix the retrieval or the prompt that caused it. Anyone selling monitoring as a cure for wrong answers is selling you a thermometer as a treatment.

Logs are personal data

Traces of real conversations contain whatever your users typed. Retention periods, redaction and access control are part of the build, not a later concern, and they are the reason this work sits next to governance rather than apart from it.

A dashboard nobody opens is not observability

The measure is whether the right person hears about the right problem without going to look. That means alerting on a small number of things people will act on, rather than forty charts that make an impressive screenshot.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

Questions

What people ask before they start

Something not covered here? Ask us directly and an engineer will answer.

Ask a question (https://codetrappers.com/contact/)

Only partly. Your existing monitoring tells you the request succeeded and took 900 milliseconds. It cannot tell you that the answer cited the wrong policy document, that the refusal rate doubled after a prompt change, or that one customer accounts for forty percent of your token spend. We instrument the AI-specific layer and send it into the tooling you already use rather than adding another dashboard.

Governance is what a system may reach and who approved it, decided before launch. Observability is what it actually did and what it cost, recorded after launch. Most teams need both, and the audit trail is where the two meet, so we usually build them to share one log rather than two. The before-launch half is an AI security and governance (https://codetrappers.com/ai-security-governance/) review.

It can, which is why retention, redaction and access control are part of the build rather than a later concern. Traces of real conversations contain whatever your users typed. We agree what is kept, for how long, and who can read it, before anything starts recording.

A set of representative tasks with known good outcomes, scored automatically on every prompt edit, model upgrade and deployment. It turns quality into a number in your pipeline instead of an impression after a demo, and it is the only reliable way to notice that an upgrade made something worse.

Start here

Ready to get started with AI Monitoring and Observability?

Tell us what you need. We'll come back with a plan and a fixed price, usually within a day.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

- A free audit first

We'll look at your site and tell you straight what needs doing.

- Fixed scope, fixed price

Agreed in writing before any work starts. No hidden fees.

- A real engineer on the call

You talk to the people who do the work, not a sales script.