Source: https://codetrappers.com/ai-search-knowledge-assistants/

Answers from your own content, with the sources attached

Retrieval-augmented search over your own pages, handbooks, tickets and product data. It cites its sources and says when it does not know. Try the one running on this site.

Retrieval-augmented search over your pages, handbooks, tickets and product data. It shows what it used to answer, and it says plainly when it does not know rather than filling the gap with something plausible.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

- 300+ migrations, zero data loss
- Core Web Vitals in the green
- 24/7 emergency cover

What you get

Inside AI Search and Knowledge Assistants

Get in touch (https://codetrappers.com/contact/)

Index what you actually have

Pages, posts, PDFs, help articles, past tickets, product records. The indexing step is where most of the quality comes from, and it is the step most implementations rush.

Retrieve before you generate

The model only ever answers from passages pulled out of your own material. That is what stops it drawing on half-remembered training data and presenting it as your policy.

Show the working

Every answer carries the pages it came from. A reader who does not trust it can check in one click, which is the difference between a tool people use and a demo they abandon.

Refuse rather than guess

The failure that damages trust is not a wrong answer, it is a confident one. Ours are built to say when the material does not cover the question, and to hand over to a person when that is the honest response.

Measure it with a real test set

A set of questions with known good answers, run on every change. Without one, tuning retrieval is guesswork and every improvement is anecdotal.

The service

What we build

Assistants for your public website

Answering from your own pages, in your own voice, with a handoff to a human at the point where a sale or a complaint starts.

Internal knowledge base assistants

For the questions your team asks each other every week. Onboarding, process, policy, what happened on that account last year.

Enterprise search across systems

One question, answered from several places at once: the intranet, the ticket history, the document store, the product catalogue.

Search over product and catalogue data

Natural language against attributes, stock, compatibility and specification, for catalogues too large to browse.

Evaluation and tuning

Existing assistant that answers badly? Often the cause is the index rather than the model, and that is measurable and fixable.

Retrieval built to be replaced in part

WordPress core embeddings are being built inside the AI plugin and are expected after 7.1. We separate indexing from scoring so that when they land you replace one half and keep the rest.

Why CodeTrappers

Try ours before you buy yours

The assistant on this site answers from these pages. It is the same architecture we would build for you, and these are the two things it taught us.

Retrieval quality is a content problem first

Our portfolio was invisible to our own search because the words visitors use, "portfolio" and "case study", appeared nowhere in the indexed text. Adding that vocabulary fixed what no amount of tuning would have.

The dangerous answer is the confident one

An article on your site quoting somebody else's prices will be retrieved and presented as yours unless the system is told, explicitly, that it must not do that.

Ask it something now

The launcher sits at the bottom of this page. Ask it what we charge, or something we

have not written about, and watch what it does with a question it cannot answer.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

Questions

What people ask before they start

Something not covered here? Ask us directly and an engineer will answer.

Ask a question (https://codetrappers.com/contact/)

A chatbot answers from a model. This answers from your own material: it retrieves passages from your pages, handbooks, tickets or product data, and generates only from those. It shows the sources it used, so a reader who does not trust an answer can check it in one click.

It says so, and hands over to a person where that is the honest response. The failure that damages trust is not a wrong answer, it is a confident one drawn from something adjacent, so we build for refusal rather than for coverage.

Usually, and usually the cause is the index rather than the model. On our own site the portfolio was invisible to search because the words visitors use, portfolio and case study, appeared nowhere in the indexed text. That is measurable and it is fixable.

No, if it is built to be replaced in part. Core embeddings are being built inside the AI plugin and are expected after 7.1, so we keep indexing separate from scoring. When they land you swap one half and keep the rest.

Start here

Ready to get started with AI Search and Knowledge Assistants?

Tell us what you need. We'll come back with a plan and a fixed price, usually within a day.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

- A free audit first

We'll look at your site and tell you straight what needs doing.

- Fixed scope, fixed price

Agreed in writing before any work starts. No hidden fees.

- A real engineer on the call

You talk to the people who do the work, not a sales script.