Source: https://codetrappers.com/custom-ai-agent-development/

Agents that do the work, not agents that describe it

Agents that do the work rather than describe it: scoped tools, permissions enforced in code, a person on anything irreversible, and monitoring for the drift that follows launch.

A chatbot answers questions. An agent looks something up, decides what to do, does it, and reports back. The difference is not the model, it is the tools it can reach and the limits around them.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

- 300+ migrations, zero data loss
- Core Web Vitals in the green
- 24/7 emergency cover

What you get

Inside Custom AI Agent Development

Get in touch (https://codetrappers.com/contact/)

Start from the job, not the model

Most requests we get are really a scheduled query nobody has time to run. That part is cheaper, faster and more reliable built without a model, and we will say so before quoting.

Give it tools rather than instructions

An agent told what to do in a prompt will improvise. An agent given a fixed set of described tools, each with a schema and a permission check, can only do the things you built.

Bound the blast radius before the first run

What it can read, what it can change, how much it can spend and how many times it can loop. Decided up front, enforced in code, not left to good behaviour.

Keep a person on anything irreversible

Sending, publishing, charging, deleting. Those get drafted and queued, because the cost of being wrong is not symmetrical and never has been.

Watch it after launch, because it will drift

Inputs change, a supplier renames a field, a model version updates. An agent that quietly started failing three weeks ago is the normal failure mode, not a dramatic one.

The service

What we build

Operations agents

The morning sweep: stuck orders, failed payments, stalled tickets, stock that has been out for a month. Read-only to begin with, which is where most of the value is anyway.

Support triage agents

Match an inbound message to the order, the shipment and the product, summarise the history, and draft a reply for a person to send.

Content and catalogue agents

Bulk work with judgement in it, such as writing descriptions from real product data and flagging the records where the data is too thin to describe honestly.

Reporting agents

Questions answered from live data rather than a dashboard nobody opens, delivered where your team already works.

Monitoring and maintenance

A retainer covering tool call volume, failure and refusal rates, cost per run, and the regression tests that catch drift before a client does.

An evaluation harness

A set of tasks with known good outcomes, run on every change. Without one, every improvement is anecdotal and every regression is a surprise.

Why CodeTrappers

Where agent projects go wrong

Three failures account for most of the disappointment, and none of them are about model quality.

Too much permission, too early

Write access granted on day one produces an agent nobody trusts enough to leave running. Read-only first earns the confidence to grant more later.

No log anyone can read

When something odd surfaces a month later, the only useful question is which tool ran with which arguments. That cannot be answered retrospectively.

Tool output treated as instruction

Text that comes back from a web page, a ticket or a document is data, not a command. An agent that does not draw that line can be steered by anyone who can write into one of those systems.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

Questions

What people ask before they start

Something not covered here? Ask us directly and an engineer will answer.

Ask a question (https://codetrappers.com/contact/)

A chatbot answers a question. An agent looks something up, decides what to do, does it and reports back. The difference is not the model, it is the tools it can reach and the limits around them.

Read-only. An agent that surfaces the five things nobody has time to check each morning is worth more in month one than a write-enabled one nobody trusts enough to leave running, and it is how a team earns the confidence to grant more.

The blast radius is decided before the first run and enforced in code: what it can read, what it can change, how much it can spend and how many times it can loop. Anything irreversible is drafted and queued for a person.

It drifts. Inputs change, a supplier renames a field, a model version updates. We watch tool call volume, failure and refusal rates and cost per run, with regression tests that catch drift before a customer does. If you already have agents running and none of that is in place, it is a piece of work on its own: AI monitoring and observability (https://codetrappers.com/ai-monitoring-observability/).

Start here

Ready to get started with Custom AI Agent Development?

Tell us what you need. We'll come back with a plan and a fixed price, usually within a day.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

- A free audit first

We'll look at your site and tell you straight what needs doing.

- Fixed scope, fixed price

Agreed in writing before any work starts. No hidden fees.

- A real engineer on the call

You talk to the people who do the work, not a sales script.