Source: https://codetrappers.com/private-self-hosted-ai/

AI that runs where your data already lives

Open-weight models on your own infrastructure, with retrieval kept inside the boundary too. Includes an honest account of what self-hosting costs and when not to do it.

For teams who cannot send client files, patient records or unreleased product data to a third-party API. Open-weight models on your own infrastructure, or in your own cloud tenancy, with the retrieval layer local too.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

- 300+ migrations, zero data loss
- Core Web Vitals in the green
- 24/7 emergency cover

What you get

Inside Private and Self-Hosted AI

Get in touch (https://codetrappers.com/contact/)

Work out what genuinely cannot leave

Usually it is less than people assume, and the answer changes the budget by an order of magnitude. Some workloads can use a hosted model with a zero-retention agreement; only the rest need hardware.

Match the model to the hardware, not the headline

An open-weight model that fits your GPU and answers your actual questions beats a larger one that needs a cluster and still has to be evaluated.

Keep retrieval local as well

Self-hosting the model and then sending your documents to a hosted embedding API defeats the purpose. The index belongs on the same side of the boundary.

Prove it against a real test set

The question is never whether a private model works, it is whether it works well enough for your task. That is measurable, and we measure it before you buy hardware.

Plan for the upgrade you will want in a year

Open weights improve quickly. A deployment that cannot swap the model without rebuilding everything around it ages badly.

The service

What we build

On-premise and private cloud deployment

In your own datacentre, your own VPC or a dedicated tenancy, with the network boundary drawn where your policy says it should be.

Model selection and evaluation

A shortlist scored against your own tasks and your own hardware, with the numbers written down rather than asserted.

Private retrieval

Index, embeddings and search all inside the boundary, so no document content crosses it at any point.

Data residency for EU clients

Processing kept in region, with the transfer question answered in a way that survives a procurement review.

Hybrid routing

Sensitive work to the private model, everything else to a hosted one, with the routing rule written in code rather than left to a person to remember.

Honest cost modelling

Hardware, hosting, operations time and the cost of the API calls you are avoiding, compared over three years rather than one month.

Why CodeTrappers

What self-hosting actually costs

We would rather you heard this now than after the purchase order, so here is the case against.

There is a quality gap, and it is task-dependent

For retrieval and extraction over your own documents it is often small enough not to matter. For open-ended reasoning it is usually still real. Which of those you are doing decides whether this is a good idea.

Somebody has to operate it

Patching, monitoring, capacity, and being awake when it stops. That is a running cost, and it is the one most often left out of the comparison.

The treadmill does not stop

A model chosen today will be behind in a year. Budget for revisiting it rather than treating the deployment as finished.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/

yusuf-mudagal-6abk5h/30min)

Questions

What people ask before they start

Something not covered here? Ask us directly and an engineer will answer.

Ask a question (https://codetrappers.com/contact/)

Often less than people assume, and the answer changes the budget by an order of magnitude. Some workloads are fine on a hosted model with a zero-retention agreement. We work out which of your data genuinely cannot leave before anyone buys hardware.

For retrieval and extraction over your own documents the gap is often small enough not to matter. For open-ended reasoning it is usually still real. Which of those you are doing decides whether this is a good idea, and we measure it against your own tasks rather than a benchmark.

Hardware or dedicated hosting, plus the operations time to patch it, monitor it and be awake when it stops. We model that against the API spend you are avoiding, over three years rather than one month, before recommending anything.

Only if the retrieval layer stays inside too. Self-hosting a model and then sending your documents to a hosted embedding API defeats the point, so the index and the embeddings sit on the same side of the boundary as the model.

Start here

Ready to get started with Private and Self-Hosted AI?

Tell us what you need. We'll come back with a plan and a fixed price, usually within a day.

Get in touch (https://codetrappers.com/contact/) Book a consultation (https://cal.com/yusuf-mudagal-6abk5h/30min)

- A free audit first

We'll look at your site and tell you straight what needs doing.

- Fixed scope, fixed price

Agreed in writing before any work starts. No hidden fees.

- A real engineer on the call

You talk to the people who do the work, not a sales script.